

# BPINN-Eikonal-Inverse-UQ

Uncertainty quantification of conduction velocities  
from noisy activation maps using  
Bayesian Physics-Informed Neural Networks

Daniele Ceccarelli

Supervisor: Prof. Manzoni

Co-supervisor: Dr. Pagani

ADVANCED PROGRAMMING FOR SCIENTIFIC COMPUTING  
Mathematical Engineering, Politecnico di Milano

A.A. 2019/2020

# Inverse UQ problem

## Inverse UQ problem

Starting from some scattered and noisy measurements  $\mathbf{D}$  reconstruct the posterior Probability Density Function of unknown parameters  $\theta$  of a given partial differential equation.

## Bayes Rule

$$\mathcal{P}(\theta|\mathbf{D}) = \frac{\mathcal{P}(\mathbf{D}|\theta)\mathcal{P}(\theta)}{\mathcal{P}(\mathbf{D})}$$

# New solution

Limitations: Classical methods involving PDE models require to solve multiple times ( $O(10^5/10^6)$ ) the PDE to sample the posterior distribution.

Our solution: **Bayesian PINN**  $\rightarrow$  a method to directly approximate the posterior distribution of a physics problem with Neural Networks.

# Physics-Informed Neural Networks (PINN)

PINN is a deterministic approach to solve forward or parameter estimation problems.

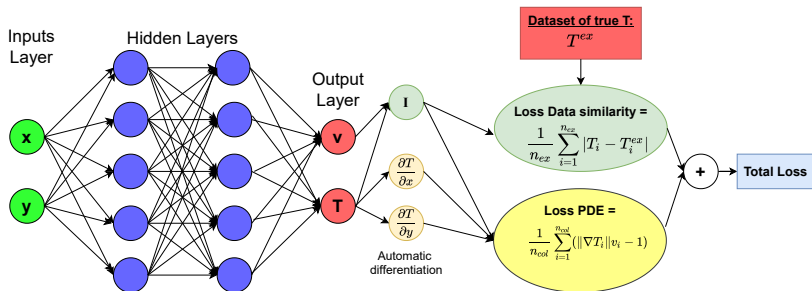


Figure 1: Example of PINN Architecture

# Bayesian PINN

Extension of PINN to a Bayesian framework in order to reconstruct the posterior distribution of our parameters.

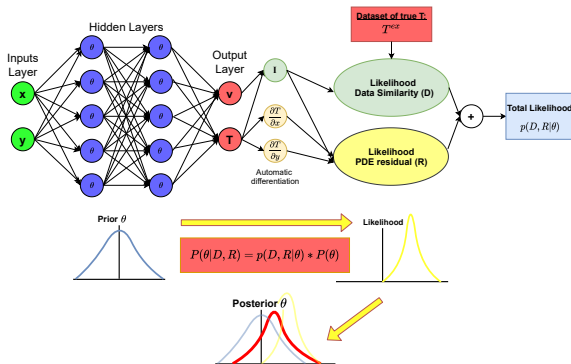


Figure 2: BPINN Architecture

# Datasets

To evaluate the Bayes formula in a NN we need 2 datasets:

- ***D***: Dataset of measurements of  $T$ ,  $\dim=n_{ex}$ , affected by Gaussian noise:

$$\hat{T}_i = T_i + \epsilon_i, \quad \epsilon_i \sim \mathcal{N}(0, \sigma^2) \quad \forall i = 1, \dots, n_{ex}.$$

- ***R***: Dataset of PDE residuals,  $\dim=n_{coll}$ , collocations points where we enforce the validity of PDE.

# Bayesian model: Likelihood functions

Let  $\theta$  be the vector that contains all the weights in the NN.

**Likelihood function:**

$$\mathcal{P}(\mathbf{D}^\theta | \theta, \sigma_D) \sim \mathcal{N}(\mathbf{D}, \Sigma_D)$$

where  $\Sigma_D = \sigma_D^2 \mathbb{I}$ , and  $\sigma_D$  can be a fixed value or an hyperparameter.

$$\mathcal{P}(\mathbf{R}^\theta | \theta, \sigma_R) \sim \mathcal{N}(\mathbf{R}, \Sigma_R)$$

where  $\Sigma_R = \sigma_R^2 \mathbb{I}$ , and also in this case  $\sigma_R$  can be fixed or an hyperparameter.

# Bayesian model: Priors

We have to specify a **prior** distribution on parameters vector  $\theta$ , for instance a t-student:

$$\mathcal{P}(\theta_j) \sim \text{Student}T(\mu^*, \lambda^*, \nu^*) \quad \forall j = 1, \dots, \dim \theta$$

If we consider both  $\sigma_D$  and  $\sigma_R$  trainable, we need a prior distribution on them, for instance an Inverse Gamma:

$$\mathcal{P}(\sigma_D^2) \sim \text{Inv} - \text{Gamma}(\alpha_1, \beta_1)$$

$$\mathcal{P}(\sigma_R^2) \sim \text{Inv} - \text{Gamma}(\alpha_2, \beta_2).$$



# Bayesian Methods

The **posterior** distribution can be computed as:

$$\begin{aligned}\mathcal{P}(\boldsymbol{\theta}, \sigma_D, \sigma_R | \mathbf{D}^\theta, \mathbf{R}^\theta) &\propto \mathcal{P}(\mathbf{D}^\theta | \boldsymbol{\theta}, \sigma_D) \mathcal{P}(\mathbf{R}^\theta | \boldsymbol{\theta}, \sigma_R) \mathcal{P}(\boldsymbol{\theta}) \mathcal{P}(\sigma_D) \mathcal{P}(\sigma_R) \\ &\propto \mathcal{N}(\mathbf{D}^\theta; \mathbf{D}, \Sigma_D) \mathcal{N}(\mathbf{R}^\theta; \mathbf{0}, \Sigma_R) \text{Student}T(\boldsymbol{\theta}; \boldsymbol{\mu}^*, \boldsymbol{\lambda}^*, \nu^*) \\ &\quad IG(\sigma_D^2; \alpha_1, \beta_1) IG(\sigma_R^2; \alpha_2, \beta_2)\end{aligned}$$

To sample from this posterior we use two different methods:

- **HMC**: Hamiltonian Monte Carlo, a classical MCMC method,
- **SVGD**: Stein Variation Gradient Descent.

## HMC

**Goal:** update weights to approximate the posterior distribution.

- Initialize  $\theta^{t_0}$ , fix  $N$ ,  $M$ ,  $L$  and  $dt$ ;
- for every  $k$  in  $1, \dots, N$ :
  - 1 Sample  $r^{t_{k-1}} \sim \mathcal{N}(0, \mathbb{I})$
  - 2  $(\theta_0, r_0) = (\theta^{t_{k-1}}, r^{t_{k-1}})$
  - 3 for  $i$  in  $0, \dots, (L-1)$ :

$$r_i = r_i - \frac{dt}{2} \nabla U(\theta_i)$$

$$\theta_{i+1} = \theta_i + dt r_i$$

$$r_{i+1} = r_i - \frac{dt}{2} \nabla U(\theta_{i+1})$$

- 4 sample  $p \sim \text{Uniform}(0, 1)$
- 5 compute  $\alpha = \min(1, \exp(H(\theta_L, r_L) - H(\theta^{t_{k-1}}, r^{t_{k-1}})))$
- 6 If  $p \geq \alpha$ , then  $\theta^{t_k} = \theta_L$ , else  $\theta^{t_k} = \theta^{t_{k-1}}$

and finally compute all the statistics we need using  $\{\theta^{t_i}\}_{i=N-M+1}^N$ .

# Application

## Inverse UQ problem

Starting from some scattered and noisy measurements  $\{\hat{T}_i\}_{i=1}^{n_{ex}}$  reconstruct the posterior Probability Density Function of conductivity tensor  $\mathbf{M}(\mathbf{x})$  using equation:

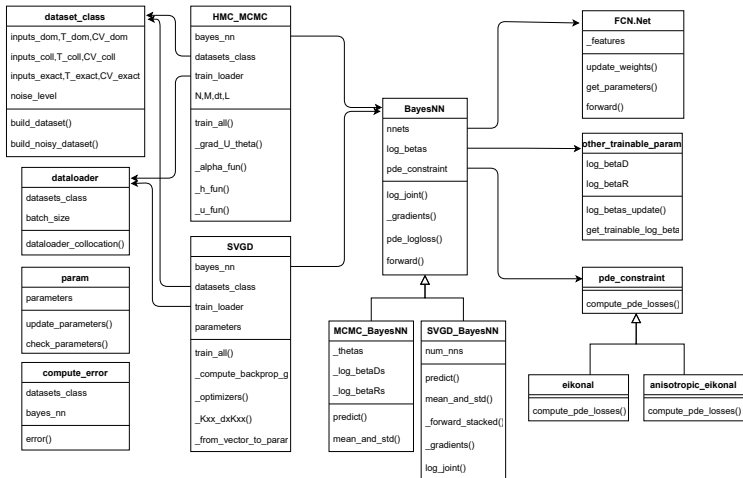
$$\begin{cases} \sqrt{\nabla T(\mathbf{x})^T \mathbf{M}(\mathbf{x}) \nabla T(\mathbf{x})} = 1 & \text{in } \Omega, \\ T(\hat{\mathbf{x}}_i) = 0 \quad \forall i = 1, \dots, N_S. \end{cases} \quad (1)$$

If  $\mathbf{M}(\mathbf{x}) = v(\mathbf{x})^2 \mathbb{I}_d$ , for a suitable conduction velocity  $v : \mathbb{R}^d \rightarrow \mathbb{R}$ , equation (1) becomes:

$$\begin{cases} \|\nabla T(\mathbf{x})\| = \frac{1}{v(\mathbf{x})} & \text{in } \Omega, \\ T(\hat{\mathbf{x}}_i) = 0 \quad \forall i = 1, \dots, N_S. \end{cases} \quad (2)$$

# Code organization

**Goal:** flexible implementation of Bayesian PINN




# External libraries

This implementation is based on **Python** (3.8.5)



## External libraries:

- Tensorflow 2
- Tensorflow-probability
- Numpy
- Matplotlib



# Fully Connected Networks

## Attribute:

**self.\_features** : a neural network implemented adding `tf.keras.layers.Dense` layers in a `tf.keras.Sequential()` model. All the layers have a *Swish* activation function and (`glorot_uniform,zeros`) initializer for (weights,bias) parameters.

## Methods:

- **forward(self, x)**: return the forward pass of input `x` in the network (`self._features(x)`);
- **get\_parameters(self)**: return all the trainable parameters in the NN in a list;
- **update\_weights(self, param)**: update the weights in all the layers.

# Bayesian PINN

## Attributes:

- **self.nnets** : neural network (or a list of nns for SVGD) to be trained (parameters vector  $\theta$ );
- **self.log\_betas** :  $\log \frac{1}{\sigma_D^2}$  and  $\log \frac{1}{\sigma_R^2}$ ;
- **self.pde\_constraint** : compute the specific PDE residual.

## Methods:

- **log\_joint(self, output, target)**: compute the log likelihood of data similarity and log prior of  $\theta$  (and  $\sigma_D$  if trainable);
- **\_gradients(self, inputs)**: compute the gradients wrt inputs (using Automatic Differentiation);
- **pde\_logloss(self, inputs)**: compute the log loss of PDE constraint (and log prior of  $\sigma_R$  if trainable);
- **predict(self, inputs)** and **mean\_and\_std(self, inputs)**: prediction, mean and std on inputs.

# HMC implementation

## Attributes:

- **self.bayes\_nn**: the BPINN object we want to train with HMC;
- **self.dataset\_class** and **self.train\_loader** : datasets and data loader classes;
- **parameters** of HMC.

## Methods:

- **train\_all(self)**: method called from main to start the training;
- **\_u\_fun(self,...)**: compute  $U(\theta) = -(\log(P(\mathbf{D}|\theta)) + \log(P(\mathbf{R}|\theta)) + \log(P(\theta)))$ ;
- **\_Grad\_U\_theta(self,...)**: compute gradient of  $U(\theta)$  wrt  $\theta$ ;
- **\_alpha\_fun(self,...)**: compute  $\alpha$  for accept/reject step.



# 1D Isotropic

1D Isotropic example:


$$\begin{cases} \left| \frac{\partial T(x)}{\partial x} \right| = \frac{1}{v(x)}, & \text{in } \Omega = (0, 1), \\ T(0) = 0. \end{cases} \quad (3)$$

with the following solutions:

$$\begin{aligned} T(x) &= 1 - e^{-2x} \\ v(x) &= \frac{1}{2} e^{2x}. \end{aligned} \quad (4)$$

Inverse UQ varying the number of exact data (10,20 and 40) and noise level on it (0.01,0.05 and 0.10)

# 1D Isotropic



**Figure 3:** Activation times and conduction velocity in 1D Isotropic example with 10 data, noise level 0.05

# 1D Isotropic




Figure 4: Comparison with 10 data and 0.01/0.05/0.10 noise level, samples from posterior distribution

# 2D Anisotropic

$$M(x) = \begin{bmatrix} a(x, y) & -c(x, y) \\ -c(x, y) & b(x, y) \end{bmatrix}, \quad \begin{cases} b(x, y) = 2a(x, y) \\ c(x, y) = a(x, y) \end{cases} \quad (5)$$




Figure 5:  $T(x,y)$  (true, mean and std of posterior distribution) in 2D Anisotropic example

# 2D Anisotropic




Figure 6:  $T(x,y)$  and  $a(x,y)$  on the line  $y=x$




Figure 7:  $a(x,y)$  (true, mean and std) in 2D Anisotropic example

# Prolate 3D

## Inverse UQ problem

Starting from some measurements of  $T$   $\{\hat{T}_i\}_{i=1}^{n_{ex}}$ , compute the posterior PDF of  $\theta$ :

$$\mathcal{P}(\theta|D, R)$$

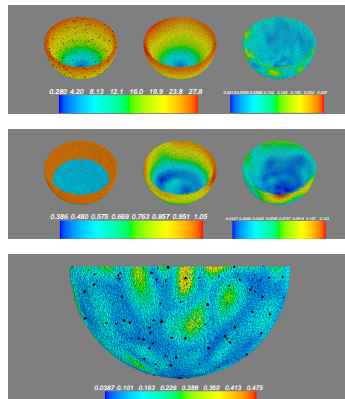



Figure 8: Activation times and conduction velocity (true, mean and std of post. distr.) in 3D prolate example

# Graph Execution


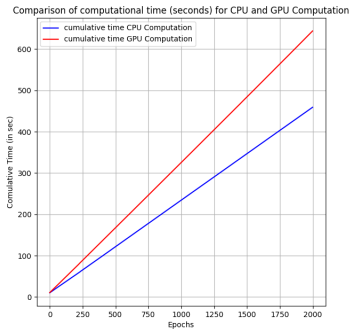
Tensorflow 2 has the possibility of use both Graph or Eager execution.

- **Eager execution:**  
Natural control flow and easy to debug;
- **Graph execution:**  
Instantiate a Graph for computation, difficult to debug but faster.



**Figure 9:** Computational time comparison for Eager and Graph execution

# GPU accelerator



**Figure 10:** Computational time comparison for CPU and GPU in a small network (3/20 architecture) and a big network (5/100 architecture)







# Conclusions

- BPINN gives an efficient and accurate inverse problem solution of Eikonal equation;
- Inverse UQ problems: preliminary results are capturing variance trend but we need to investigate more (the relation with prior distributions and likelihood functions);
- We test the model in different numerical experiments (up to 3D Isotropic and 2D Anisotropic);
- Flexible implementation, could be used with minor changes also to solve Inverse UQ for other PDEs.

## Further developments

- Explore the relationship with **different Bayesian models** (priors and likelihoods);
- Implement different MCMC methods (for instance **NUTS**);
- **Active learning** algorithm to dynamically minimize the uncertainty;
- **Transfer learning** with subdomains;
- Anisotropy and **more realistic** scenarios.

# References

-  M. Raissi and P. Perdikaris and G.E. Karniadakis, “Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations”.
-  Liu Yang and Xuhui Meng and George Em Karniadakis, “B-PINNs: Bayesian physics-informed neural networks for forward and inverse PDE problems with noisy data”.
-  Sahli Costabal, Francisco and Yang, Yibo and Perdikaris, Paris and Hurtado, Daniel E. and Kuhl, Ellen, “Physics-informed neural networks for cardiac activation mapping”.
-  Sun, Luning and Wang, Jian-Xun, “Physics-constrained Bayesian neural network for fluid flow reconstruction with sparse and noisy data”

# Appendix: SVGD

- Initialize  $N$  neural networks weights  $\theta^i \forall i = 1, \dots, N$ ;
- For every epochs  $e = 1, \dots, E$ :
  - 1 Compute the log posterior  $L(\theta^i) \forall i = 1 \dots, N$
  - 2 Compute the gradients  $\nabla_{\theta^i} L(\theta^i)$  by back-propagation
  - 3 Compute

$$\phi(\theta^i) = \frac{1}{N} \sum_{j=1}^N [k(\theta^i, \theta^j) \nabla_{\theta^i} L(\theta^i) + \nabla_{\theta^i} k(\theta^i, \theta^j)] \quad \forall i = 1 \dots, N \quad (6)$$

- 4 Update  $\theta^i = \theta^i + \epsilon \phi(\theta^i)$  or use a Stochastic Gradient Descent method,  $\forall i = 1, \dots, N$ .

# Appendix: Posterior PDF

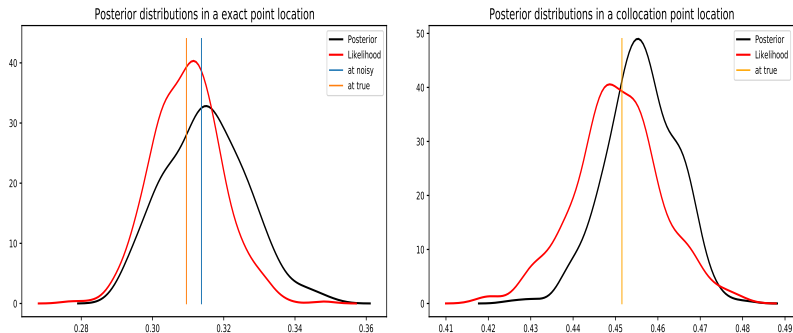




Figure 11: Posterior PDF and Likelihood on Activation times locations

# Appendix: Influence of Eikonal



**Figure 12:** Comparison of activation times with 10 exact data,  $\text{noise}_{lv} = 0.05$ , “With Eikonal” on the left, “Without Eikonal” on the right

# Appendix: Posterior of $\sigma_D$

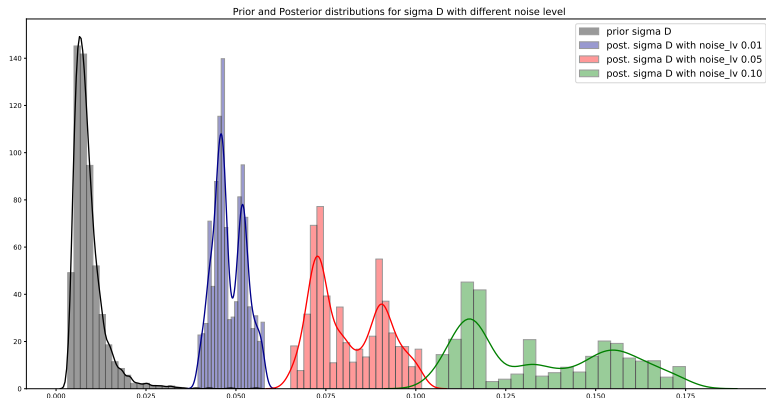


Figure 13:  $\sigma_D$  prior and posteriors for noise level = 0.01, 0.05, 0.10